

Natural Language Processing

Lecture 2: Basic Text Preprocessing

Salima Lamsiyah

University Luxembourg, FSTM, DCS, MINE Research Group

Salima.lamsiyah@uni.lu



Revisions

1. What is NLP ?
2. Give some examples of NLP applications
3. Why NLP is hard?
4. What are the different levels of ambiguity ?
5. How to deal with ambiguity?
6. What are the main components of NLP?

Lecture Plan

1. Corpora and Words
2. Text Normalization
 - a. Tokenization
 - b. Stemming
 - c. Lemmatization
3. Sentence Segmentation
4. Regular Expressions

Corpora and Words

Corpora

- A corpus is a collection of text\speech
 - Often annotated in some way
 - Sometimes just lots of text
- A text is produced by
 - a specific writer(s),
 - at a specific time,
 - in a specific variety,
 - of a specific language,
 - for a specific function.

Corpora vary along dimension like

- **Language:** 7097 languages in the world
- **Variety,** like African\Mainstream American Language varieties.
- **Code switching,** e.g., Spanish/English, Hindi/English:
- **Genre:** newswire, fiction, scientific articles, Wikipedia
- **Author Demographics:** writer's age, gender, ethnicity, socioeconomic class , ...
- **Time:** language changes over time

Corpus datasheets | data statement

Motivation:

- Why was the corpus collected?
- By whom?
- Who funded it?

Situation: In what situation was the text written?

Collection process: If it is a subsample how was it sampled? Was there consent? Pre-processing?

- +Annotation process, language variety, demographics, etc.

Gebru et al (2020), Bender and Friedman (2018)

How many words in a sentence?

- Example: "They lay back on the San Francisco grass and looked at the stars and their dreams took flight"
- **Tokens** are the total number of running words.
- **Types** are the number of distinct words in a corpus.
- How Many:
 - 18 tokens (or 17)
 - 16 types (or 15)

How many words in a sentence?

- I'm
- **Orthographically** one word (in the English writing system)
- But **grammatically** two words:
 1. the subject pronoun I
 2. the verb 'm, short for am.

How many words in a corpus?

- N = number of tokens
- V = vocabulary = set of types, $|V|$ is size of the vocabulary.
- **Heaps'Law** = $|V| = kN^\beta$ where often $\beta = 0.5$, $10 < k < 100$
- **Vocab size for a text goes up with the square root of its length in words**

| Corpus | Tokens = N | Types = $ V $ |
|---------------------------------|--------------|---------------|
| Switchboard phone conversations | 2.4 million | |
| Shakespeare | 884,000 | |
| COCA | 440 million | |
| Google N-grams | 1 trillion | |

How many words in a corpus?

- N = number of tokens
- V = vocabulary = set of types, $|V|$ is size of the vocabulary.
- **Heaps'Law** = $|V| = kN^\beta$ where often $\beta = 0.5$, $10 < k < 100$
- **Vocab size for a text goes up with the square root of its length in words**

| Corpus | Tokens = N | Types = $ V $ |
|---------------------------------|--------------|---------------|
| Switchboard phone conversations | 2.4 million | 20 thousand |
| Shakespeare | 884,000 | 31 thousand |
| COCA | 440 million | 2 million |
| Google N-grams | 1 trillion | 13+ million |

Text Normalization

Text Normalization

- Every NLP application requires **Text Normalization**:
 1. Word Tokenization
 2. Word Normalization (normalizing word formats)
 - Stemming
 - Lemmatization
 - Lowercasing
 3. Sentence Segmentation

Text Tokenization

Word Tokenization: Space-based tokenization

- **Word Tokenization** consists of splitting a sentence into an ordered list of individual words, usually referred to as tokens.
- A very simple way for tokenization **Space-based tokenization**
 - For languages that use space characters between words
 - Examples: Arabic, Latin, Greek, ...
- Unix Tools for space-based tokenization
 - The “**tr**” command
 - [Inspired by Ken Church’s Unix for Poets](#)
 - Given a text file, output the word tokens and their frequencies

Tokenization in languages without spaces

- Many languages (like Chinese, Japanese, Thai) don't use spaces to separate words!
- How do we decide where the token boundaries should be?

How to choose tokens in Chinese

- Chinese words are composed of characters called "hanzi" (汉字) (or sometimes just "zi")
- Each one represents a meaning unit called a morpheme.
- Each word has on average 2.4 of them.
- But deciding what counts as a word is complex and not agreed upon.

How to do choose tokens in Chinese?

- 姚明进入总决赛 “Yao Ming reaches the finals”
• yáo míng jìn rù zǒng jué sài

• 3 words?

- 姚明 进入 总决赛
- YaoMing reaches finals

Chinese Treebank

• 5 words?

- 姚 明 进 入 总 决 赛
- Yao Ming reaches overall finals

Peking University

• 7 words?

- 姚 明 进 入 总 决 赛
- Yao Ming enter enter overall decision game

Just use characters

Tokenization across languages

- So Chinese uses characters (zi) as tokens
 - But that doesn't work for, e.g., Thai and Japanese
 - These differences make it hard to use words as tokens
- And there's another reason why we don't use words as tokens!

Morphological Typology

- Dimensions along which languages vary
- Two are salient for tokenization:
 1. number of morphemes per word
 2. how easy it is to segment the morphemes

Words have parts

- **Morpheme:** a minimal meaning-bearing unit in a language.
 - **fox:** one morpheme
 - **cats:** two morphemes **cat** and **-s**
- **Morphology:** the study of morphemes

Types of morphemes

- **root:** central morpheme of the word; supplying the main meaning
- **affix:** adding additional meanings

Examples:

worked

root work

ffix -ed

glasses

root glass

affix -es

Types of affixes

- **Inflectional morphemes**
 - grammatical morphemes
 - often syntactic role like agreement
 - ed past tense on verbs
 - s/-es plural on nouns
- **Derivational morphemes**
 - more idiosyncratic in application and meaning
 - often change grammatical class
 - care* (noun)
 - + -full → *careful* (*adjective*)
 - + -ly → *carefully* (*adverb*)

Issues in Tokenization

- **Can't just blindly remove punctuation:**
 - m.p.h., Ph.D.,
 - prices ([\\$45.55](#))
 - dates ([20/12/1992](#))
 - URLs (<http://www.stanford.edu>)
 - hashtags ([#nlproc](#))
 - email addresses (someone@cs.colorado.edu)
- **Critic: a word that doesn't stand on its own**
 - "are" in [we're](#), French "je" in [j'ai](#), "le" in [l'honneur](#)
- **When should multiword expressions (MWE) be words?**
 - [New York](#), [rock 'n' roll](#)

Word Tokenization: Subword Tokenization

- Instead of
 - white-space segmentation
 - single-character segmentation
- **Use the data** to tell us how to tokenize.
- **Subword tokenization** (because tokens can be parts of words as well as whole words)

Subword Tokenization Methods: Examples

- Three common algorithms:
 - **WordPiece Tokenizer** (Schuster and Nakajima, 2012)
 - **Byte-Pair Encoding (BPE) Tokenizer** (Sennrich et al., 2016)
 - **SentencePiece Tokenizer** (Kudo, 2018)
- All have 2 parts:
 - A token **learner** that takes a raw training corpus and induces a vocabulary (a set of unique tokens).
 - A token **segmenter** that takes a raw test sentence and tokenizes it according to that vocabulary

Subword Tokenization: Byte Pair Encoding

- **Byte Pair Encoding** is a subword tokenization method that splits words into smaller parts by merging the most frequent pairs of characters or subwords.
- **How It Works:**
 - It starts by treating each character in a word as its own token.
 - The most frequent pair of characters is merged into a single token.
 - This process repeats until a desired vocabulary size is reached.
- **Example:** Let's tokenize "**lowering**" using BPE:
 - Step 1: Start with individual characters: l o w e r i n g
 - Step 2: Merge frequent pairs:
 - "low" + "er" + "ing"
 - Final tokens: low, er, ing
- So, "**lowering**" becomes: **low er ing**

Subword Tokenization: WordPiece Tokenizer

- **WordPiece** starts with individual characters and learns which subword combinations to merge based on maximizing the likelihood of the training data.
- **How It Works:**
 - Words are split into the smallest units or subwords.
 - These subwords are merged based on their likelihood in the data, and if a word is not in the vocabulary, it will be split into smaller known parts.
- **Example:** For "playing":
 - WordPiece might split it into: play, ##ing
 - The ## symbol indicates that "ing" is a suffix attached to another part of the word.
- So, "playing" becomes: **play ##ing**

Subword Tokenization: SentencePiece Tokenizer

- SentencePiece is a tokenization algorithm that doesn't require space as a boundary for words. It treats the entire sentence as a stream of characters and breaks it into subwords or tokens based on frequency.
- How It Works:
 - It works on raw text (without spaces as word boundaries).
 - It segments the text into subword units using a probabilistic model.
- Example: Let's use "I am learning" with SentencePiece:
 - SentencePiece could split it into: I, _am, _learn, ing
 - The _ symbol before "am" and "learn" represents a space that was part of the original sentence.
- So, "I am learning" becomes: I _am _learn ing

Subword Tokenization Methods

Check this [link](#) for more details and examples

Tokenization in NLTK

✓ 0s

```
▶ from nltk.tokenize import word_tokenize  
text = "I am learning NLP."  
tokens = word_tokenize(text)  
print(tokens)
```

```
⇒ ['I', 'am', 'learning', 'NLP', '.']
```

Treebank Tokenizer

```
▶ from nltk.tokenize import TreebankWordTokenizer  
  
tokenizer = TreebankWordTokenizer()  
sentence = "They'll meet at 5:00 p.m. on Jan. 5th, 2021."  
tokens = tokenizer.tokenize(sentence)  
print(tokens)
```

```
⇒ ['They', "'ll", 'meet', 'at', '5:00', 'p.m.', 'on', 'Jan.', '5th', ',', '2021', '.']
```

Regular Expression Tokenization

✓ 0s

```
▶ from nltk.tokenize import regexp_tokenize  
text = "The price is $45.50 for a book."  
tokens = regexp_tokenize(text, r'\w+|\$[\d\.\.]+')  
print(tokens)
```

```
⇒ ['The', 'price', 'is', '$45.50', 'for', 'a', 'book']
```

Tokenization in SpaCy

```
' is ➔ import spacy  
  
# Load the English language model  
nlp = spacy.load('en_core_web_sm')  
  
text = "I'm learning NLP with spaCy!"  
doc = nlp(text)  
  
# Tokenizing the text  
tokens = [token.text for token in doc]  
print(tokens)
```

```
→ ['I', "'", 'm', 'learning', 'NLP', 'with', 'spaCy', '!']
```

Word Normalization

Word Normalization

- Putting words/tokens in a standard format
 - U.S.A. or USA
 - making or made
 - Fed or fed
 - am, is, be, are
 - studies, studying, studied, study
 - ...

Word Normalization: Case Folding

- Applications like Information Retrieval: reduce all letters to lower case is helpful
- For Named Entity Recognition, Machine Translation, Part-of-speech tagging, and other applications
 - Case is not helpful (**US** versus **us** is important)

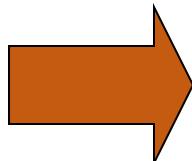
Word Normalization: Stemming

- **Stemming** is the process of chopping off affixes (prefix, suffix, infix) crudely from a word in order to obtain a word **stem (root)**.
- [Porter Stemmer](#) is the most used stemming method, Developed by Martin Porter in the 1980s
 - Based on a series rules
 - Some sample rules:
 - **ATIONAL => ATE** (e.g., communicational => communicate)
 - **ING => remove it if stem contains vowel** (e.g., playing => play)
 - **SSES => SS** (e.g., processes => process)

Word Normalization: Stemming

- Example

This was not the map we found in Billy Bones's chest, but an accurate copy, complete in all things-names and heights and soundings-with the single exception of the red crosses and the written notes.



Thi wa not the map we found in Billi Bone s chest but an accur copi complet in all thing name and height and sound with the singl except of the red cross and the written note .

Word Normalization: Lemmatization

Lemmatization returns the canonical form, dictionary form of a word. The output we will get after lemmatization is called ‘lemma’.

- Am, are, is => **be**
- Cat, cats, cat's, cats' => **cat**
- I was in a meeting => **I be in a meeting**
- We were meeting constantly => **We be meet constantly**

Sentence Segmentation

- **Sentence segmentation** consists of splitting the text into a set of sentences using a specific sentence splitter.
- !, ? mostly unambiguous but period “.” is very ambiguous because:
 - Sentence boundary
 - Abbreviations like Inc. or Dr.
 - Numbers like .02% or 4.3
- Common algorithm: Tokenize first: use rules or Machine Learning to classify a period as either (a) part of the word or (b) a sentence-boundary.
 - An abbreviation dictionary can help
- Sentence segmentation can then often be done by rules based on this tokenization

Text Preprocessing Libraries

1. Natural Language Toolkit (NLTK): <https://www.nltk.org/>
2. spaCy Library: <https://spacy.io/>

<https://www.activestate.com/blog/natural-language-processing-nltk-vs-spacy/>

Other Text Preprocessing Techniques

- Stop-words removal
- Special Character Removal
- Sometimes Removing Emails, URLs, XML tags, ...

<https://kavita-ganesan.com/text-preprocessing-tutorial/#.YzBavuxBy3K>

Regular Expressions

Regular Expressions

- Regular expression (RE) is a formal language that specifies text search strings.
- Formally, a *regular expression* is an algebraic notation for characterizing a set of strings.
- Very useful for searching in texts, when we have a ***pattern*** to search for and a ***corpus*** of texts to search through.
- How can we search for any of these?
 - woodchuck
 - woodchucks
 - Woodchuck
 - Woodchucks



Regular Expressions: Concatenation

- Concatenation

| Reglaur Expression | Example Patterns Matched |
|--------------------|---|
| /woodchuck/ | "groundhogs, also known as woodchuck , are rodent of the family Sciuridae" |
| /a/ | "Bring me a cup of tea" |
| /!/ | " Just enjoy the NLP course! " |

Regular Expressions: Disjunction

- Disjunction

The use of the brackets [] to specify a disjunction of characters.

The use of the brackets [] Plus the dash - to specify a range.

| Reglaur Expression | Match |
|--------------------|------------------------|
| /[wW]oodchuck/ | woodchuck or Woodchuck |
| /[abc]/ | 'a', 'b', or, 'c' |
| /[1234567890]/ | Any digit |

| Reglaur Expression | Match |
|--------------------|----------------------|
| /[A-Z]/ | An upper case letter |
| /[a-z]/ | A lower case letter |
| /[0-9]/ | A single digit |

Regular Expressions: Negation in Disjunction

- Negations $[^a]$:
 - Caret \wedge means negation only when it is the first symbol in []

| Reglaur Expression | Match (single character) |
|--------------------|--------------------------|
| $/[^A-Z]/$ | Not an upper case letter |
| $/[^aA]/$ | Neither 'a' nor 'A' |
| $/[^e^]/$ | Neither 'e' nor '^' |
| $/[e^]/$ | Either 'e' or '^' |
| $/a^b/$ | The pattern a^b |

Regular Expressions: ? * + .

- The question mark **?** means the preceding character or nothing.
- The asterisk ***** indicates zero or more occurrences of the preceding element.
- The plus sign **+** indicates one or more occurrences of the preceding element.
- The period **(./.)**, a wildcard expression that matches any single character except a carriage return character \r.

| Reglaur Expression | Match |
|------------------------|---|
| /woodchucks ? / | woodchuck or woodchucks |
| /ab * c/ | ac, abc, abbc, abbcc, and so on. |
| /ab + c/ | abc, abbc, abbcc, and so on, but not “ac” |
| /beg . n/ | Begin, began, beg’n, and so on. |

Regular Expressions: Anchors \wedge $\$$

- Anchors are special characters that don't match any characters but anchor regular expressions to particular places in a string.
- The most common anchors are the caret \wedge and the dollar $\$$
- The caret \wedge matches the start of the line.
- **Remark:** The caret has 3 uses:
 - To match a start of a line,
 - To indicate a negation inside of square brackets,
 - And just mean a caret.
- The dollar sign $\$$ matches the end of a line.

Regular Expressions: Anchors ^ \$

- Examples:

| Reglaur Expression | Match | Example Pattern Matched |
|--------------------|---------------------------|-------------------------------------|
| /^A-Z]/ | Start of the line | Salima Lamsiyah |
| /^A-Za-z]/ | Start of the line | <u>1.</u> “HELLO” |
| \.\$ | End of line | The end. <u>_</u> |
| .\$ | End of line | The end? <u>_</u> The end! <u>_</u> |
| /^The dog\.\$/ | Start and end of the line | The dog. |
| \b | Word boundary | \bthe\b/ returns the |

Regular Expressions: More Disjunction

- Woodchuck is another name for groundhog!
 - The pipe | for disjunction

| Reglaur Expression | Match |
|--------------------|---|
| | Either the string woodchuck or groundhog |
| | Either the string guppy or guppies |
| | Woodchuck, woodchuck, Groundhog, groundhog |

Regular Expressions: More Disjunction

- Woodchuck is another name for groundhog!
 - The pipe | for disjunction

| Reglaur Expression | Match |
|---------------------------|--|
| /woodchuck groundhog / | Either the string woodchuck or groundhog |
| /gupp (y ies)/ | Either the string guppy or guppies |
| [gG]roundhog [Ww]oodchuck | Woodchuck, woodchuck, Groundhog, groundhog |

Regular Expressions

- Regular Expression Operator precedence hierarchy

| | |
|-----------------------|---------------|
| Parenthese | () |
| Counters | * + ? {} |
| Sequences and anchors | The ^my end\$ |
| Disjunction | |

More Regular Expressions Operators

| RE | Expansion | Match | First Matches |
|----|--------------|-----------------------------|---------------|
| \d | [0-9] | any digit | Party_of_5 |
| \D | [^0-9] | any non-digit | Blue_moon |
| \w | [a-zA-Z0-9_] | any alphanumeric/underscore | Daiyu |
| \W | [^\w] | a non-alphanumeric | !!! |
| \s | [\r\t\n\f] | whitespace (space, tab) | |
| \S | [^\s] | Non-whitespace | in_Concord |

| RE | Match |
|-------|---|
| * | zero or more occurrences of the previous char or expression |
| + | one or more occurrences of the previous char or expression |
| ? | exactly zero or one occurrence of the previous char or expression |
| {n} | n occurrences of the previous char or expression |
| {n,m} | from n to m occurrences of the previous char or expression |
| {n,} | at least n occurrences of the previous char or expression |
| {,m} | up to m occurrences of the previous char or expression |

Example

- Find all the instances of the word "there" in a text.
- You can use RegEx in Python
- <https://regex101.com/>
- https://www.w3schools.com/python/python_regex.asp

Example

- Find me all the instances of the word "there" in a text.
 - `/there/` => Misses capitalized examples
 - `/[Tt]here/` => Incorrectly returns *therefore*
 - `\b[Tt]here\b/` => will not return (`there_` or `there10`)
 - `/[^a-zA-Z][Tt]here[^a-zA-Z]/` => Might not return there when it begins a line
 - Final Regular Expression: `/(^|[^a-zA-Z])[Tt]here([^a-zA-Z])|$/`

Error

- The process we just went through was based on fixing two kinds of errors:
 1. Matching strings that we should not have matched (therefore, weathered, feathered)
False positives (Type I errors)
 2. Not matching string that we should have matched (There)
False negatives (Type II errors)

Error

- In NLP we are always dealing with these kinds of errors.
- Reducing the error rate for an application often involves two antagonistic efforts:
 - **Increasing accuracy or precision** (minimizing false positives)
 - **Increasing coverage or recall** (minimizing false negatives).

Exercice

- Write regular expressions for the following languages.
 1. The set of all alphabetic strings;
 2. The set of all lower case alphabetic strings ending by b;
 3. All strings that start at the beginning of the line with an integer and that end at the end of the line with a word;
 4. Check date in DD/MM/YYYY and DD.MM.YYYY formats
 5. Write a regular expression that matches Column 1 Column 2 Column 3

You can use <https://regex101.com/>

More Regular Expressions (Optional)

Substitutions

- Substitution in Python and UNIX commands:

s/regexp1/pattern/

e.g. **s/colour/color/**

More Regular Expressions

Capture Groups

- Say we want to put angles around all numbers:
e.g. the 40 boxes => the <40> boxes
- Use () to "capture" a pattern into a numbered register (1, 2, 3...)
- Use \1 to refer to the contents of the register

S / ([0-9]+) /<\1>/

More Regular Expressions

Capture Groups: Multiple Registers

/the (.*)er they (.*), the \1er we \2/

- Matches:
the **faster** they **ran**, the **faster** we **ran**
- But not
the **faster** they **ran**, the **faster** we **played**

More Regular Expressions

But suppose we don't want to capture?

- Parentheses have a double function: grouping terms, and capturing
- Non-capturing groups: add a ?: after parentheses:

/(:?some|a few) (people|cats) like some \1/

- matches
 - some cats like some cats

but not

- some cats like some people

Simple Regular Expressions Application: ELIZA

Chatbot ELIZA

- ELIZA an early chatbot that imitated a Rogerian psychotherapist
 - *Joseph Weizenbaum, 1966.*
- Uses pattern matching to match, e.g.,:
 - “I need X”
- and translates them into, e.g.
 - “What would it mean to you if you got X?”

Simple Regular Expressions Application: ELIZA

Chatbot ELIZA

- Let's chat with ELIZA <https://web.njit.edu/~ronkowitz/eliza.html>.

User₁: Men are all alike.

ELIZA₁: IN WHAT WAY

User₂: They're always bugging us about something or other.

ELIZA₂: CAN YOU THINK OF A SPECIFIC EXAMPLE

User₃: Well, my boyfriend made me come here.

ELIZA₃: YOUR BOYFRIEND MADE YOU COME HERE

User₄: He says I'm depressed much of the time.

ELIZA₄: I AM SORRY TO HEAR YOU ARE DEPRESSED

a sample conversation with ELIZA in 1966

Simple Regular Expressions Application: ELIZA

How ELIZA Works

```
s/.* I'M (depressed|sad) .*/I AM SORRY TO HEAR YOU ARE \1/  
s/.* I AM (depressed|sad) .*/WHY DO YOU THINK YOU ARE \1/  
s/.* all .*/IN WHAT WAY/  
s/.* always .*/CAN YOU THINK OF A SPECIFIC EXAMPLE/
```

Summary

- Regular expressions play a surprisingly large role
 - Sophisticated sequences of regular expressions are often the first model for any text processing application
- For hard tasks, we use machine learning classifiers
 - But regular expressions are still used for pre-processing, or as features in the classifiers
 - Can be very useful in capturing generalizations

Reading

- Please refer to this document : [Chapter 2](#)

- Regular Expressions:

<https://docs.python.org/3/library/re.html#>

https://en.wikipedia.org/wiki/Regular_expression

<https://docs.python.org/3/library/re.html#>

<https://regex101.com/>

<https://regexr.com/>

https://www.youtube.com/watch?v=V_BozMwoYe4&ab_channel=techTFQ

<https://www.w3resource.com/python-exercises/re/>

- ELIZA Chatbot

<https://web.njit.edu/~ronkowit/eliza.html>

<http://psych.fullerton.edu/mbirnbaum/psych101/eliza.htm>

Next Class

- Edit Minimum Distance
- Text Representation Part 1

